

OpenMP offloading with Fortran

Soner Steiner

Intel Instructor

VSC Research Center, TU-WIEN

soner dot steiner at tuwien dot ac dot at
sonersteiner at gmail dot com

Intel oneAPI HPC Workshop- Agenda

Online 19th-20th September 2023

DAY 2 – THEME: PROGRAMMING WITH ONEAPI

Time	Session name / description	Presenter
09:30	Welcome	<i>Soner/Claudia</i>
09:40	Porting CUDA code to SYCL using the Compatibility Tool	<i>Georg</i>
10:30	HANDS-ON WITH COMPATIBILITY TOOL A hands-on lab session where you can try porting a CUDA code to oneAPI with the help of the Compatibility Tool	<i>Georg</i>
11:00	Coffee	
11:15	OFFLOADING WITH C/C++ and FORTRAN Offloading using OpenMP mainly in C/C++ Offloading using OpenMP in FORTRAN Automatic offloading using DO CONCURRENT	<i>Soner</i>
12:00	Lunch	
13:00	LAB3: HANDS-ON OFFLOADING WITH OPENMP	<i>Soner</i>
14:30	Coffee	
14:45	LAB4: HANDS-ON VTUNE A hands-on lab session where you can use the Vtune and Advisor profilers to assess the performance of some example codes.	<i>Soner</i>
16:00	End of Day 2	

Agenda

- Automatic offloading with DO CONCURRENT
- Offloading with OpenMP
- OpenMP Target Construct
- Managing Device Data
- Implementing Parallelism

DO CONCURRENT

- ♦ *Embarrassingly parallel problems:*
 - Can be distributed across several cores with little to no effort.
 - E.g.: large number of independent data records, graphics rendering, ...
 - DO CONCURRENT can be used!
- ♦ *Non embarrassingly parallel problems:*
 - Any parallel problem with interdependency between processes, requires communication and synchronization.

DO CONCURRENT

- ♦ What is **do concurrent**?

It is a promise from the programmer to compiler that the code inside loop can be safely parallelized.

- ♦ What **do concurrent** is not?

It is not a guarantee that the loop will run in parallel, maybe the compiler determines that a serial execution would be more efficient.

DO CONCURRENT

do [,] concurrent ([type-spec ::] index-spec-list [, scalar-mask-expr])

- `type-spec` (if present) specifies the type and kind of the index variables
- `index-spec-list` is a list of index specifications
do concurrent (i=1:n, j=1:m)

DO CONCURRENT, Examples

```
do concurrent( i = 1:n )  
  c(i) = a(i) + b(i)  
end do
```

```
do concurrent(i = 1:grid_size)  
  h(i) = exp(-decay * (i - icenter)**2)  
end do
```

Demo at the end, ... [demos/00_demo/](#)

OpenMP Offload Constructs

Be aware:

Offloading to a device with oneAPI icx, icpx and ifx works only with Intel GPUs at the moment.

We are expecting it in the coming year

Just like the backend SW we have used with SYCL for the nvidia GPUs .

OpenMP Offload Constructs

Device Code

```
omp target [clause[[,clause]...] structured-block  
omp declare target [function-definitions-or-declarations]  
omp declare target [variable-definitions-or-declarations]
```

Worksharing

```
omp teams [clause[[,clause]...] structured-block  
omp distribute [clause[[,clause]...] for-loops
```

Memory operations

```
map ([[map-type-modifier[,map-type-modifier] map-type:] list) map-type := alloc |  
tofrom | to | from | release | delete map-type-modifier := always  
omp target data clause[[[,clause]...] structured-block  
omp target enter data clause[[[,clause]...]  
omp target exit data clause[[[,clause]...]  
omp target update clause[[[,clause]...]
```

OpenMP Offload Language

C/C++	Fortran
#pragma omp target <i>[clause[[,]clause]...]</i> <i>structured-block</i>	!\$omp target <i>[clause[[,]clause]...]</i> <i>structured-block</i> !\$omp end target
#pragma omp target data <i>[clause[[,]clause]...]</i> <i>structured-block</i>	!\$omp target <i>[clause[[,]clause]...]</i> <i>structured-block</i> !\$omp end target data
#pragma omp teams <i>[clause[[,]clause]...]</i> <i>structured-block</i>	!\$omp teams <i>[clause[[,]clause]...]</i> <i>structured-block</i>
#pragma omp distribute <i>[clause[[,]clause]...]</i> <i>structured-block</i>	!\$omp distribute <i>[clause[[,]clause]</i> <i>...]</i> <i>structured-block</i>

OpenMP 4.0 for Devices - Constructs

#pragma omp target **!\$omp target**

- **target** construct transfer control and data from the host to the device

- Syntax (C/C++)

```
#pragma omp target [clause[,] clause],...]  
structured-block
```

- Syntax (Fortran)

```
!$omp target [clause[,] clause],...]  
structured-block  
!$omp end target
```

- Clauses

```
device(scalar-integer-expression)  
map([{alloc | to | from | tofrom}:] list)  
if(scalar-expr)
```

Target construct

target [clause]

- Offloads a code region to a target device
- Sequential and synchronous by default

clause : device, private,
firstprivate, map, allocate
Sync: nowait, depend

```
1 program target_construct_1
2   implicit none
3
4   integer :: i
5   integer :: a(100), b(100), c(100)
6
7   do i=1, 100
8     a(i) = 3
9     b(i) = 7
10  end do
11
12  !$omp target
13    do i=1, 100
14      c(i) = a(i) + b(i)
15    end do
16  !$omp end target
17
18  do i=1, 100
19    write(*,*), a(i), b(i), c(i)
20  end do
21
22 end program target_construct_1
```

Device
code

Device Clause

```
!$omp target device(i)
```

- Specify which device to offload to in a multi-device environment
 - Device number an integer
 - Assignment is implementation-specific
 - Usually start at 0 and sequentially increments
 - Works with **target**, **target data**, **target enter/exit data**, **target update** directives

Device Clause

Target device construct
Integer numbers, 0, 1, 2
specify the device

```
:>$ sycl-ls  
[opencl:acc:0] Intel(R) FPGA Emulation Platform for OpenCL(TM), Intel(R) FPGA  
[opencl:cpu:1] Intel(R) OpenCL, Intel(R) Core(TM) i9-10900K CPU @ 3.70GHz 3.0  
[opencl:gpu:2] Intel(R) OpenCL HD Graphics, Intel(R) UHD Graphics 630 [0x9bc5]
```

```
ifx -qopenmp -fopenmp-targets=spir64 -o target_construct_7.x target_construct_7.f90  
:>$ OMP_TARGET_OFFLOAD="MANDATORY"  
:>$ LIBOMPTARGET_PLUGIN=LEVEL0  
:>$ LIBOMPTARGET_DEBUG=1
```

```

program target_construct_2
  use omp_lib
  implicit none

  integer :: i
  integer, parameter :: n=1000000
  integer :: a(n), b(n), c(n)

  do i=1, n
    a(i) = 3
    b(i) = 7
  end do

  !$omp target device(2)
    do i=1, n
      c(i) = a(i) + b(i)
    end do
  !$omp end target

  do i=1, n
    write(*,*), a(i), b(i), c(i)
  end do

end program target_construct_2

```

Device
code

Target device construct

target device

- Specify which device to offload to in a multi-device environment
- Device number is an integer, usually starts at 0 and sequentially increments
- Works with **target**, **target data**, **target enter \ exit data**, **target update** directives

OpenMP* Device Parallelism

```
5  double precision :: start_time, end_time
6  integer :: i
7  integer, parameter :: n=100000000
8  integer :: a(n), b(n), c(n)
9
10 do i=1, n
11     a(i) = 3
12     b(i) = 7
13 end do
14
15 start_time = omp_get_wtime()
16 !$omp target device(2)
17     !$omp parallel do
18         do i=1, n
19             c(i) = a(i) + b(i)
20         end do
21     !$omp end parallel do
22 !$omp end target
23 end_time = omp_get_wtime()
24
25 do i=1, n
26     write(*,*), a(i), b(i), c(i)
27 end do
28 write(*,*) "Work took: ", end_time - start_time, " seconds"
```

**Device
code**

target [clause]

- Offloads a code region to a target device
- Sequential and synchronous by default

Why NOT parallel for?

- CPU parallelism differs from GPU – shared memory systems
- omp parallel for threads will use only 1 Streaming Multiprocessor (SM) to synchronize
- Need a different level of parallelism to step over multiple SM

OpenMP* Device Parallelism

```
5  double precision :: start_time, end_time
6  integer :: i
7  integer, parameter :: n=100000000
8  integer :: a(n), b(n), c(n)
9
10 do i=1, n
11     a(i) = 3
12     b(i) = 7
13 end do
14
15 start_time = omp_get_wtime()
16 !$omp target device(2)
17 !$omp teams distribute parallel do
18     do i=1, n
19         c(i) = a(i) + b(i)
20     end do
21 !$omp end teams distribute parallel do
22 !$omp end target
23 end_time = omp_get_wtime()
24
25 ! do i=1, n
26 !     write(*,*), a(i), b(i), c(i)
27 ! end do
28 write(*,*) "Work took: ", end_time - start_time, " seconds"
```

**Device
code**

target [clause]

- Offloads a code region to a target device
- Sequential by default

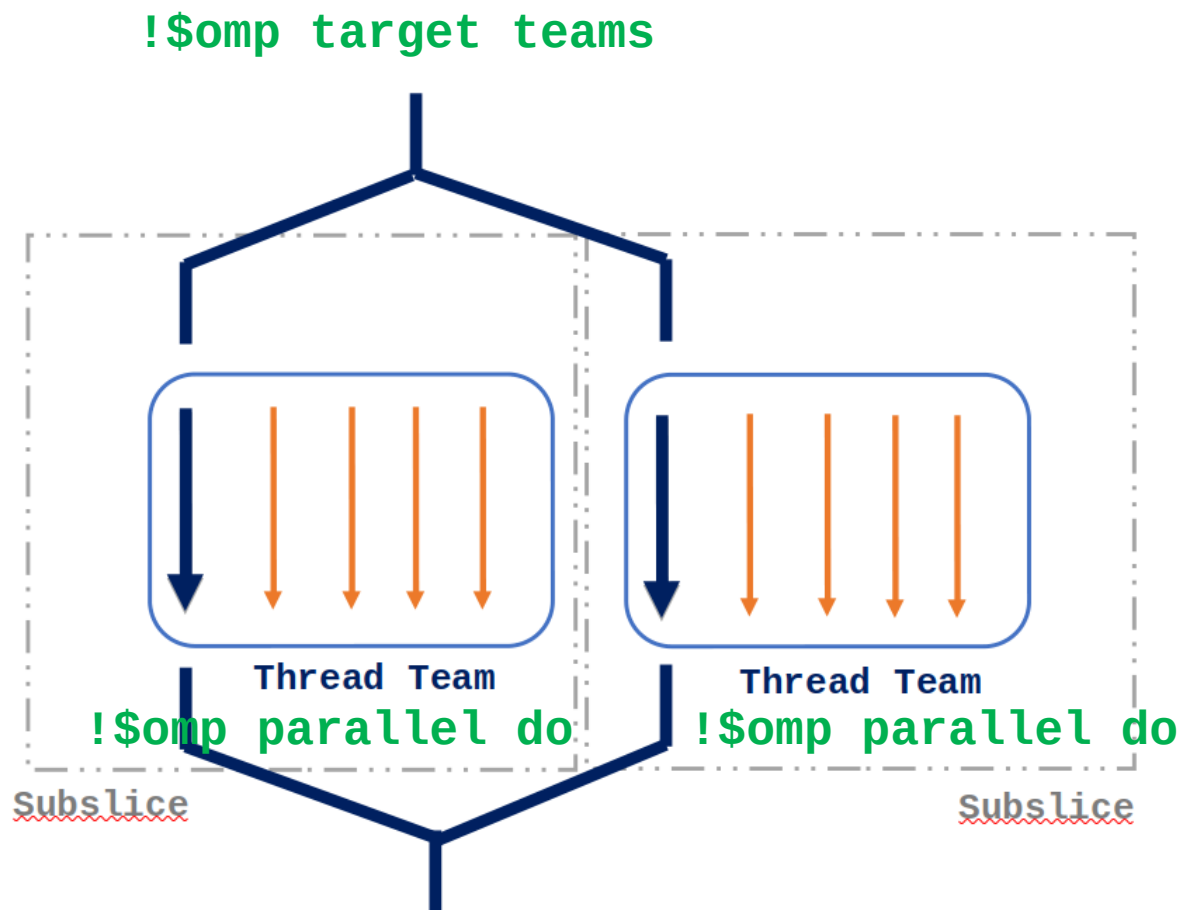
Target teams

- creates a league of teams where the primary thread of each team executes the teams region

2.13.21 target teams Construct

Teams Construct

OpenMP	GPU Hardware
SIMD	SIMD Lane (Channel)
Thread	SIMD Thread mapped to an EU
Team	Group of threads mapped to a Subslice
League	Multiple Teams mapped to a GPU



OpenMP* Worksharing

```
5  double precision :: start_time, end_time
6  integer :: i
7  integer, parameter :: n=100000000
8  integer :: a(n), b(n), c(n)
9
10 do i=1, n
11     a(i) = 3
12     b(i) = 7
13 end do
14
15 start_time = omp_get_wtime()
16 !$omp target teams distribute parallel do
17     do i=1, n
18         c(i) = a(i) + b(i)
19     end do
20 !$omp end target teams distribute parallel do
21 end_time = omp_get_wtime()
22
23 ! do i=1, n
24 !     write(*,*), a(i), b(i), c(i)
25 ! end do
26 write(*,*) "Work took: ", end_time - start_time, " seconds"
27
28 end program target_construct_5
29
30
```

**Device
code**

target teams distribute

shortcut for specifying a target construct containing a teams distribute construct and no other statements.

**target teams distribute
parallel do**

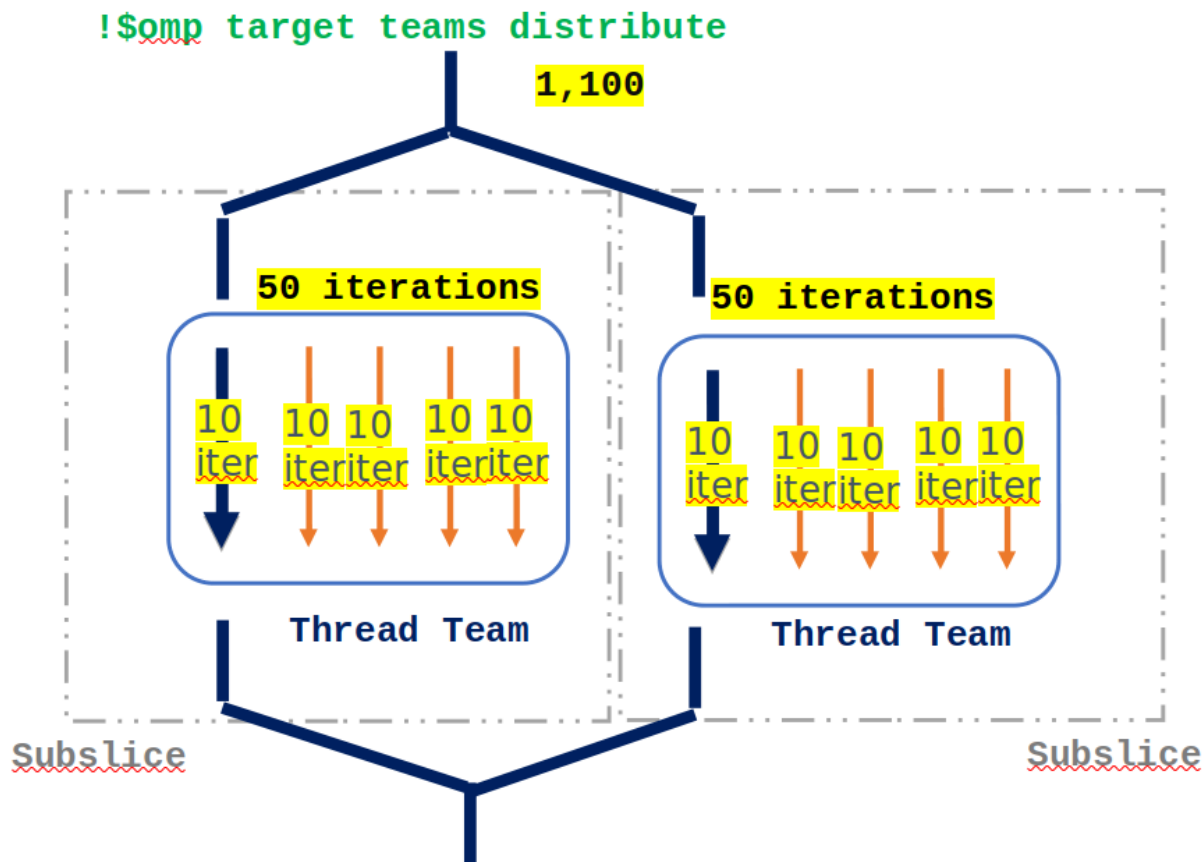
parallel worksharing-loop construct is a shortcut for specifying a target construct containing a teams distribute parallel worksharing-loop construct and no other statements

[2.13.11 teams distribute Construct](#)

[2.13.25 Target Teams Distribute Parallel Worksharing-Loop Construct](#)

Teams Distribute Construct

```
!$omp target teams  
distribute parallel do  
  do k=1,100  
    c(k) = a(k) + b(k)  
  end do  
!$omp end target teams  
distribute parallel do
```



Calling functions inside Target region

```
subroutine devicefunc()  
!$omp declare target  
device_type(device)  
...  
end subroutine  
  
program main  
!$omp target  
    call devicefunc()  
!$omp end target  
end program
```

declare target

compiles a version of the function/subroutine for the target device

Function compiled for both host execution and target execution by default

**device_type(host |
nohost | any)**

2.12.7 declare target Directive

Managing Device Data

Data Offload

- The host and devices have separate memory spaces.
 - When code is offloaded, data needs to be mapped.
 - By default, copied to the device on entry and from the device to the host on exit.
-
- Data environment is **lexically scoped**
 - Data environment is destroyed at closing curly brace
 - Allocated buffers/data are automatically released

Default Behaviour:

The compiler identifies variables that are used in the target region.

All accessed arrays are copied from host to device and back

Copying a and b back is not necessary: it was not changed.

```
5  double precision :: start_time, end_time
6  integer :: i
7  integer, parameter :: n=100000000
8  integer :: a(n), b(n), c(n)
9
10 do i=1, n
11     a(i) = 3
12     b(i) = 7
13 end do
14
15 start_time = omp_get_wtime()
16 !$omp target device(2)
17 !$omp parallel do
18     do i=1, n
19         c(i) = a(i) + b(i)
20     end do
21 !$omp end parallel do
22 !$omp end target
23 end_time = omp_get_wtime()
24
25 do i=1, n
26     write(*,*), a(i), b(i), c(i)
27 end do
28 write(*,*) "Work took: ", end_time - start_time, " seconds"
```


Target map construct, map clause

- Use **map** clause to manually determine how an original variable in a data environment is mapped to a corresponding variable in a device data environment
 - omp target **map** (*map-type: list*)
 - Available map-type
 - **to** : alloc and assign value of original variable on target region entry
 - **from** : alloc and assign value to original variable on target region exit
 - **tofrom**: default, both to and from

Using MAP to refine behaviour:

```
5  double precision :: start_time, end_time
6  integer :: i
7  integer, parameter :: n=1000
8  integer :: a(n), b(n), c(n)
9
10 do i=1, n
11     a(i) = 3
12     b(i) = 7
13 end do
14
15 start_time = omp_get_wtime()
16 !$omp target map(to:a) map(to:b) map(tofrom:c)
17     !$omp parallel do
18         do i=1, n
19             c(i) = a(i) + b(i)
20         end do
21     !$omp end parallel do
22 !$omp end target
23 end_time = omp_get_wtime()
24
25 do i=1, n
26     write(*,*) c(i)
27 end do
28 write(*,*) "Work took: ", end_time - start_time, " seconds"
```

Unnecessary to copy
a and b back to the host

OpenMP Offload Constructs

▪ Device Code

- **omp target** [*clause*[[,*clause*]]...] *structured-block*
 - **omp declare target** [*function-definitions-or-declarations*]
 - **omp declare target** [*variable-definitions-or-declarations*]
- ## ▪ Worksharing
- **omp teams** [*clause*[[,*clause*]]...] *structured-block*
 - **omp distribute** [*clause*[[,*clause*]]...] *for-loops*

▪ Memory operations

- **map** ([[*map-type-modifier*[[,*map-type*]]] *list*)
map-type := **alloc** | **tofrom** | **to** | **from** | **release** | **delete** *map-type-modifier* := **always**
- **omp target data** *clause*[[[,*clause*]]...] *structured-block*
- **omp target enter data** *clause*[[[,*clause*]]...]
- **omp target exit data** *clause*[[[,*clause*]]...] *structured-block*
- **omp target update** *clause*[[[,*clause*]]...] *structured-block*

Compiling and running on your laptop

```
$ ifx -fopenmp -fopenmp-targets=spir64 -O2 -g  
-o 00-hello_world_omp.x 00-hello_world_omp.f90
```

build

```
$ ./00-hello_world_omp.x
```

run

Setting the number
of threads used via
the shell variable

```
$ env OMP_NUM_THREADS=8 ./00-Hello.x
```

Using the
gcc compiler

```
$ g++ -fopenmp -O2 -g -o 00-Hello.x 00-Hello.cpp
```

Or using a simple makefile

```
1 # Compiler settings
2 FXX = ifx
3 #FXXFLAGS = -qopenmp -fopenmp-targets=spir64 -qopt-report
4 FXXFLAGS = -qopenmp -fopenmp-targets=spir64 -O2 -g
5
6 # Setting the source and binary files
7 SRC = $(wildcard *.f90)
8 BIN = $(SRC:.f90=)
9
10 # Rules
11 default: 00_hello_world_omp
12
13 00_hello_world_omp: 00_hello_world_omp.f90
14     $(FXX) $(FXXFLAGS) -o 00_hello_world_omp.x 00_hello_world_omp.f90
15
16 clean:
17     @$(RM) $(BIN)
18
19 # Setting the independent commands
20 .PHONY: default clean
```

Makefile saved as
makefile_simple.mak

\$ make -f makefile_simple.mak

Invoke for compiling



Notices & Disclaimers

Performance varies by use, configuration, and other factors. Learn more at www.Intel.com/PerformanceIndex.

Performance results are based on testing as of dates shown in configurations and may not reflect all publicly available updates. See configuration disclosure for details.

Your costs and results may vary.

Intel technologies may require enabled hardware, software or service activation.

© Intel Corporation. Intel, the Intel logo, and other Intel marks are trademarks of Intel Corporation or its subsidiaries. Other names and brands may be claimed as the property of others.